



ELARA MVP: COMPREHENSIVE TECHNICAL AUDIT REPORT

Current Implementation vs. Required Architecture

Report Date: October 10, 2025

Assessment Type: Gap Analysis - MVP Readiness Evaluation

Scope: Full-Stack Architecture, Data Models, APIs, Services, Compliance

EXECUTIVE SUMMARY

Overall Completion: ~25-30% of MVP Requirements Met

The current Elara implementation represents an **early-stage prototype** with solid UI/UX foundations and basic AI integration. However, it falls critically short of the Product/Tech Brief specifications across architecture, data persistence, service layers, and essential features. The gap between current state and MVP requirements represents **6-10 weeks of full-stack development work**.

Critical Finding: The application lacks the fundamental infrastructure (database, microservices, authentication, automations) required for production deployment. Current implementation is suitable only as a **UI/UX demo or proof-of-concept**.

1. ARCHITECTURE GAP ANALYSIS

1.1 Required Architecture (Product Brief Section 2)

Required Backend Structure:

API Gateway → Orchestrator (LLM agents) → 8 Microservices

- └─ NLP Service (sentiment + intent classifier)
- └─ Memory/RAG Service (conversation memory, user embeddings)
- └─ Doctor Graph Service (search + ranking)
- └─ Automation Service (jobs: reminders, check-ins)
- └─ Insights Service (Mirror summaries, trends)

- └ Wearable Connector (HealthKit/Google Fit/Oura)
- └ Compliance Guard (prompt shields, policy checks)
- └ Analytics/Events Service (product analytics)

Datastores:

- └ PostgreSQL (core entities)
- └ Redis (sessions/queues)
- └ S3/Blob (reports(exports)
- └ Vector DB (memories, doc profiles, tone snippets)

1.2 Current Implementation

Actual Structure:

Single Express.js Monolith

- └ /api/classify (OpenAI classification)
- └ /api/elara/reply (OpenAI chat)
- └ /api/mock/doctors (static JSON)
- └ /api/mock/tests (static JSON)
- └ /health (health check)

Data Storage:

- └ In-Memory (MemStorage class)
- └ localStorage (client-side only)

1.3 Architecture Gap Summary

Component	Required	Current	Status
Microservices	8 services	0 services	✗ 0%
API Gateway	Yes (orchestrator)	No	✗ Missing
Service Mesh	Yes	No	✗ Missing
PostgreSQL	Yes (primary DB)	Configured but unused	⚠ 10%
Redis	Yes (sessions/queues)	No	✗ Missing
Vector DB	Yes (RAG/embeddings)	No	✗ Missing
S3/Blob Storage	Yes (reports)	No	✗ Missing

Architecture Score: 5/100 ✗

2. DATA MODEL COMPARISON

2.1 Required Schema (Product Brief Section 3)

The Product/Tech Brief specifies **12 database tables** with complex relationships:

A. User Management (2 tables)

```
users(id uuid pk, name text, locale text, tz text, created_at timestamptz  
consents(user_id uuid, type text, granted boolean, granted_at timestamptz
```

B. Conversation System (2 tables)

```
conversations(id uuid pk, user_id uuid fk, started_at timestamptz)  
messages(id uuid pk, conversation_id uuid fk,  
role text check(role in ('user','assistant','system')),  
text text, sentiment text, intent text, created_at timestamptz)
```

C. Doctor Matching (2 tables)

```
doctors(id uuid pk, name text, specialty text, clinic text, city text,  
personality_tags text[], modalities text[], rating numeric,  
lat numeric, lng numeric, meta jsonb)  
matches(id uuid pk, user_id uuid, doctor_id uuid, score numeric,  
reasons text[], created_at timestamptz)
```

D. Booking System (1 table)

```
bookings(id uuid pk, match_id uuid fk, slot timestamptz, status text,  
confirmation_ref text, created_at timestamptz)
```

E. Automation Engine (1 table)

```
jobs(id uuid pk, user_id uuid, type text, run_at timestamptz,  
payload jsonb, status text)
```

F. Wellness Tracking (3 tables)

```

daily_checkins(id uuid pk, user_id uuid, mood text, energy int,
    note text, created_at timestamptz)
weekly_reflections(id uuid pk, user_id uuid, week_start date,
    summary text, highlights jsonb)
wearable_metrics(id uuid pk, user_id uuid, day date,
    sleep_hours numeric, steps int, hrv numeric, rhr numeric, source text)

```

G. Feature Flags (1 table)

```
feature_flags(user_id uuid, key text, enabled boolean)
```

2.2 Current Implementation

```
-- ONLY TABLE DEFINED:
users(id serial pk, username text not null unique, password text not null)
```

2.3 Data Model Gap Analysis

Database Component	Required Tables	Implemented	Gap
User Management	2 tables	1 partial	⚠️ 50%
Conversations	2 tables	0	✗ 0%
Doctor Matching	2 tables	0	✗ 0%
Bookings	1 table	0	✗ 0%
Automations	1 table	0	✗ 0%
Wellness Tracking	3 tables	0	✗ 0%
Feature Flags	1 table	0	✗ 0%
Consents	1 table	0	✗ 0%

Database Schema Score: 8/100 ✗

Critical Issues:

- UUID primary keys required → Using serial integers
- User table has wrong fields (username/password vs. name/locale/tz)
- No foreign key relationships defined

- No indexes for performance
 - No migrations infrastructure
 - **91% of required schema is missing**
-

3. API ENDPOINTS GAP ANALYSIS

3.1 Required Endpoints (Product Brief Section 4)

Specified REST APIs:

1. **POST /chat/send** → { message }
 - Returns: { reply, sentiment, intent, next_actions[] }
 - **Current:** Partial (split into /classify and /reply)
 - **Gap:** No structured next_actions, no intent in response
2. **GET /doctors/search** → Query params: specialty, city, mood_tag
 - Returns: [{doctor, score, reasons[]}]
 - **Current:** ✗ Missing entirely
 - **Gap:** Using Bubble.io directly from frontend (security issue)
3. **POST /bookings** → { doctor_id, slot }
 - Returns: { booking_id, status, confirmation_ref }
 - **Current:** ✗ Missing entirely
 - **Gap:** No backend booking persistence
4. **POST /checkins** → { mood, energy, note }
 - **Current:** ✗ Missing entirely
 - **Gap:** Daily check-in feature not implemented
5. **GET /mirror/weekly** → Optional: { week_start }
 - Returns: { summary, trend }
 - **Current:** ✗ Missing entirely
 - **Gap:** Wellness reflection feature not implemented
6. **GET /wearables/summary** → Descriptive aggregates (if consented)
 - **Current:** ✗ Missing entirely
 - **Gap:** Wearable integration not started

7. POST /automations/schedule → Reminders/check-ins

- **Current:** ✗ Missing entirely
- **Gap:** No automation system

8. GET /legal/disclaimer → Latest non-medical copy for locale

- **Current:** ✗ Missing entirely
- **Gap:** No compliance infrastructure

3.2 API Implementation Status

Endpoint	Required	Implemented	Functional	Score
POST /chat/send	✓ Yes	⚠ Partial	60%	⚠
GET /doctors/search	✓ Yes	✗ No	0%	✗
POST /bookings	✓ Yes	✗ No	0%	✗
POST /checkins	✓ Yes	✗ No	0%	✗
GET /mirror/weekly	✓ Yes	✗ No	0%	✗
GET /wearables/summary	✓ Yes	✗ No	0%	✗
POST /automations/schedule	✓ Yes	✗ No	0%	✗
GET /legal/disclaimer	✓ Yes	✗ No	0%	✗

API Completeness Score: 15/100 ✗

4. MICROSERVICES ANALYSIS

4.1 Required Services (Product Brief Section 2 & 6)

Service 1: NLP Service ✗ NOT IMPLEMENTED

- **Purpose:** Sentiment + intent classification
- **Required:** Fast, cheap model for classifying latest message + 3 previous
- **Output:** { sentiment, intent, topics[] }
- **Current:** Basic OpenAI classification endpoint exists
- **Gap:** Not a separate service, no topic extraction, no context window (3 prev messages)

Service 2: Memory/RAG Service ✗ NOT IMPLEMENTED

- **Purpose:** Conversation memory, user embeddings
- **Required:** Vector DB for semantic search across conversation history
- **Current:** localStorage-based memory (client-side only)
- **Gap:** No server-side memory, no embeddings, no RAG, no vector search

Service 3: Doctor Graph Service NOT IMPLEMENTED

- **Purpose:** Doctor search + ranking algorithm
- **Required Ranking Formula:**

```
final_score = 0.35*specialty + 0.25*keywords + 0.2*personality_fit +
             0.1*distance + 0.1*quality
```

- **Current:** Bubble.io API called directly from frontend
- **Gap:** No ranking algorithm, no personality matching, no geo-distance calculation

Service 4: Automation Service NOT IMPLEMENTED

- **Purpose:** Job scheduling (reminders, check-ins)
- **Required:** BullMQ/Celery job queue
- **Jobs:** Reminder (+1h, +48h post-booking), abandoned booking follow-up
- **Current:** None
- **Gap:** No job system, no queue, no scheduled tasks

Service 5: Insights Service NOT IMPLEMENTED

- **Purpose:** Weekly "Mirror" summaries, trend analysis
- **Required:** Generate summary every Monday 9am user TZ
- **Current:** None
- **Gap:** No wellness tracking, no trend analysis, no scheduled summary generation

Service 6: Wearable Connector NOT IMPLEMENTED

- **Purpose:** HealthKit/Google Fit/Oura integration
- **Required:** Daily summary ingestion with consent checks
- **Current:** None
- **Gap:** No wearable integration, no OAuth connectors

Service 7: Compliance Guard NOT IMPLEMENTED

- **Purpose:** Prompt shields, policy checks, refusal patterns
- **Required:** Prepend non-medical policy to every LLM call
- **Current:** Basic system prompts only
- **Gap:** No compliance middleware, no audit logging, no PII/PHI handling

Service 8: Analytics/Events X NOT IMPLEMENTED

- **Purpose:** Product analytics (PostHog/Mixpanel)
- **Required KPIs:** Activation, DAU/WAU, conversion, retention, NPS/CES
- **Current:** None
- **Gap:** No analytics integration, no event tracking

4.2 Microservices Score: 0/100 X

All 8 required microservices are missing. Current architecture is a single monolithic Express server.

5. STATE MACHINE LOGIC (Section 5)

5.1 Required Flow Logic

The Product Brief specifies a **six-vertical state machine**:

```
START → EMOTION_DETECT
  | intent == education → EDUCATION_CARD
  | intent == doctor_match → MATCH_AND_RANK → CONCIERGE_FLOW
  | intent == reflection → REFLECTION_FLOW
  | intent == action → CONCIERGE_FLOW
  | else → SUGGEST_NEXT (buttons: Learn | Connect | Reflect)
```

MATCH_AND_RANK Flow:

1. Extract features: {symptom_terms, mood_tag, location, prefs}
2. Query DoctorGraph.query(features)
3. Rank by: mood_alignment + specialty + ETA
4. Return top-3 + action buttons

CONCIERGE_FLOW:

1. On 'Book': create booking, schedule reminder job (+1h, +48h follow-up)
2. On 'Maybe later': offer check-in reminder toggle

REFLECTION_FLOW:

1. Assemble last 7 days: messages + checkins + wearable summaries
2. Generate warm, non-medical weekly summary

3. Save to weekly_reflections table
4. CTA: "Set gentle focus for next week?" (sleep | hydration | calm)

EDUCATION_CARD:

1. Return short, friendly explainer
2. Add 2-3 micro-actions
3. Safe CTA: "Would you like general lifestyle tips or talk to a specialist?"

5.2 Current Implementation

Actual Logic:

- Linear flow: WelcomeScreen → ChatScreen → DoctorMatchScreen → BookingScreen
- No state machine orchestrator
- No intent-based routing
- No reflection or education flows
- Minimal concierge automation

5.3 State Machine Score: 15/100

Gaps:

-  No orchestrator pattern
-  No intent-based branching
-  No reflection flow
-  No education cards
-  No "SUGGEST_NEXT" buttons
-  Basic doctor matching exists (manual, no ranking)
-  No concierge automation (reminders, follow-ups)

6. CORE USER STORIES COMPLIANCE

User Story 1: Natural Language Chat with Mood/Intent Detection

Acceptance Criteria: Sentiment label + intent saved per message

Requirement	Status	Implementation
Chat in natural language	 Works	OpenAI integration functional

Requirement	Status	Implementation
Mood detection	⚠ Partial	Client-side tone detection (warm/soothing/bright)
Intent detection	✗ Missing	No intent classification (doctor_match/reflection/education)
Save per message	✗ Missing	No database, localStorage only

Score: 40/100 ⚠

User Story 2: Top-3 Doctor Matches + Booking

AC: 3 cards with specialty, personality tags, distance, next availability; booking returns confirmation object

Requirement	Status	Implementation
Top-3 doctor matches	⚠ Partial	Bubble API returns all matching category
Specialty shown	✓ Works	Displayed in UI
Personality tags	✗ Missing	Not shown or ranked by personality
Distance calculation	✗ Missing	No geo-distance, no user location
Next availability	✗ Missing	Hardcoded static dates
Booking confirmation object	✗ Missing	No backend persistence

Score: 25/100 ✗

User Story 3: Follow-ups & Reminders

AC: After booking, schedule reminder + 48h check-in

Requirement	Status
Reminder system	✗ Missing
48-hour check-in	✗ Missing
Job scheduling infrastructure	✗ Missing

Score: 0/100 ✗

User Story 4: Daily Check-in + Weekly "Mirror" Summaries

AC: Weekly summary generated; trend view shows mood tag counts vs. days

Requirement	Status
Daily emotional check-in	✗ Missing
Weekly Mirror summary	✗ Missing
Trend visualization	✗ Missing
Monday 9am generation (user TZ)	✗ Missing

Score: 0/100 ✗

User Story 5: Wearable Connection (Opt-in)

AC: Sleep duration/consistency + activity ingestion; descriptive UI copy only

Requirement	Status
HealthKit/Google Fit connector	✗ Missing
Consent management	✗ Missing
Daily summary ingestion	✗ Missing
Descriptive (non-medical) display	✗ Missing

Score: 0/100 ✗

7. COMPLIANCE & SAFETY ANALYSIS (Section 9)

7.1 Required Compliance Features

Feature	Required	Current	Gap
Prompt Shields	Prepend non-medical policy to every LLM call	⚠ Basic system prompt	No dedicated compliance middleware
Refusal Patterns	Specific medical boundary responses	⚠ Generic prompts	No structured refusal logic
PII/PHI Handling	Encrypt at rest, tokenized IDs, least-privilege	✗ None	No encryption, no access controls
Audit Logging	Log all tool calls (inputs/outputs, truncated)	✗ None	No audit trail

Feature	Required	Current	Gap
Content Filters	Block unsafe supplement/medication advice	✗ None	No content filtering

Compliance Score: 10/100 ✗

Critical Risks:

- No PII/PHI encryption** - Health data stored in plain localStorage
- No audit trail** - Cannot track AI recommendations for liability
- No content moderation** - AI could provide unsafe health advice
- Exposed API token** - Bubble.io token hardcoded in frontend (Line 5, DoctorMatchScreen.tsx)

8. ACCEPTANCE CRITERIA ANALYSIS (Section 11)

8.1 MVP Acceptance Criteria Status

Criterion	Target	Current	Status
Sentiment/intent accuracy	≥85% on 100-sample test	Unknown (no test set)	✗ Not measured
Doctor search p95 latency	≤1.2s	N/A (frontend call)	✗ Not implemented
Booking success rate	≥95% confirmation saved	0% (no backend)	✗ Failing
Reminder job reliability	Fire reliably (logged)	N/A (no jobs)	✗ Not implemented
Weekly Mirror generation	Every Monday 9am user TZ	N/A	✗ Not implemented
Compliance test pass rate	100% (no medical advice)	Unknown (no tests)	✗ Not measured

Acceptance Criteria Met: 0/6 ✗

9. SECURITY VULNERABILITIES

9.1 Critical Security Issues

Issue	Severity	Location	Impact

Issue	Severity	Location	Impact
Exposed API Token	● CRITICAL	DoctorMatchScreen.ts x:5	Anyone can access Bubble.io API
No Authentication	● CRITICAL	All endpoints	Backend APIs publicly accessible
No Input Validation	● HIGH	All POST endpoints	Potential injection attacks
No Rate Limiting	● MEDIUM	OpenAI endpoints	API cost exploitation
Client-side Data Storage	● MEDIUM	localStorage	Health data accessible in browser
No HTTPS Enforcement	● MEDIUM	Server config	Data transmitted insecurely
No CORS Restrictions	● MEDIUM	Express config	Cross-origin attacks possible

Security Score: 15/100 ●

10. MISSING FEATURES INVENTORY

10.1 Complete Missing Features List

Authentication & User Management:

- ✗ Magic link authentication
- ✗ OTP authentication
- ✗ JWT token management
- ✗ Session handling
- ✗ User profile CRUD
- ✗ Consent management UI
- ✗ Feature flag system

Conversation Management:

- ✗ Multiple conversation threads
- ✗ Chat inbox/history view
- ✗ Conversation archiving
- ✗ Message search
- ✗ Export conversations

Doctor Matching:

- ✗ Ranking algorithm (5-factor scoring)
- ✗ Personality tag matching
- ✗ Geo-distance calculation
- ✗ Availability checking
- ✗ Top-3 recommendation logic
- ✗ "More options" flow

Booking System:

- ✗ Backend booking API
- ✗ Booking status management
- ✗ Confirmation email/SMS
- ✗ Calendar integration
- ✗ Reminder scheduling
- ✗ Booking history

Wellness Tracking:

- ✗ Daily check-in feature
- ✗ Weekly Mirror summaries
- ✗ Mood trend visualization
- ✗ Energy tracking
- ✗ Wearable data ingestion

Automation:

- ✗ Job queue system
- ✗ Scheduled reminders
- ✗ 48h follow-ups
- ✗ Abandoned booking recovery
- ✗ n8n workflow integration

Admin/Partner Portal:

- ✗ Doctor management interface
- ✗ Booking dashboard
- ✗ Analytics dashboard
- ✗ User management tools

Compliance:

- ✗ Audit logging system
- ✗ Content moderation
- ✗ PII/PHI encryption
- ✗ Legal disclaimer API

- ✗ Privacy policy integration
-

11. TECHNOLOGY STACK COMPARISON

11.1 Recommended vs. Current

Component	Recommended (Brief)	Current	Alignment
Frontend	React Native / Next.js	React + Vite	⚠ Partial (web only)
Backend	NestJS / FastAPI	Express.js	⚠ Acceptable but monolithic
LLM	OpenAI function calling	OpenAI chat completions	⚠ Partial (no function calling)
Database	PostgreSQL + Redis	None (in-memory)	✗ Critical gap
Vector DB	pgvector / Weaviate	None	✗ Missing
Job Queue	BullMQ / Celery	None	✗ Missing
Auth	OAuth + JWT	None	✗ Missing
Analytics	PostHog / Mixpanel	None	✗ Missing

12. DEVELOPMENT ROADMAP TO MVP

Phase 1: Foundation (3-4 weeks)

Priority: CRITICAL

Week 1-2: Database & Core Backend

- Implement complete PostgreSQL schema (12 tables)
- Set up Drizzle migrations
- Build User management APIs
- Build Conversation management APIs
- Implement proper error handling

Week 3-4: Authentication & Security

- Magic link authentication
- JWT token system
- Session management (Redis)
- Move API keys to backend

- Implement rate limiting
 - Add input validation (Zod)
-

Phase 2: Core Features (3-4 weeks)

Priority: HIGH

Week 5-6: Booking System

- Booking CRUD APIs
- Status management workflow
- Confirmation email integration
- Real-time availability checking
- Booking history endpoint

Week 7-8: Doctor Matching Service

- Implement 5-factor ranking algorithm
 - Personality tag matching
 - Geo-distance calculation
 - Top-3 recommendation logic
 - "More options" filtering
-

Phase 3: Advanced Features (2-3 weeks)

Priority: MEDIUM

Week 9-10: Automation & Wellness

- Job queue setup (BullMQ)
- Reminder scheduling system
- 48h follow-up automation
- Daily check-in feature
- Weekly Mirror summary generation

Week 11: Polish & Compliance

- Audit logging system
 - Compliance middleware
 - Content moderation
 - Analytics integration (PostHog)
 - Admin dashboard basics
-

13. COST & EFFORT ESTIMATION

13.1 Development Effort

Refer to the loe

13.2 Infrastructure Costs (Monthly)

Service	Provider	Estimated Cost
PostgreSQL	Neon / Supabase	\$25-50
Redis	Upstash / Redis Cloud	\$10-30
Vector DB	Weaviate Cloud	\$25-100
S3 Storage	AWS S3	\$5-20
OpenAI API	OpenAI	\$100-300 (usage-based)
Analytics	PostHog	\$0-50
Email/SMS	SendGrid/Twilio	\$20-100
TOTAL		\$185-550/month

14. RISK ASSESSMENT

14.1 Critical Risks

Risk	Probability	Impact	Mitigation
Data Loss	HIGH	CRITICAL	Implement database immediately
Security Breach	MEDIUM	CRITICAL	Add authentication + encryption
Compliance Violation	HIGH	HIGH	Implement audit logging + compliance guard
Scalability Issues	HIGH	MEDIUM	Migrate to microservices architecture
AI Cost Overrun	MEDIUM	MEDIUM	Add rate limiting + caching

15. FINAL RECOMMENDATIONS

15.1 Immediate Actions (Next 2 Weeks)

1. **STOP** considering this for mobile deployment in current state
2. **PRIORITIZE** database implementation (PostgreSQL schema)
3. **SECURE** the Bubble.io API token (move to backend)
4. **IMPLEMENT** authentication system (magic link)
5. **BUILD** booking backend APIs with persistence

15.2 Go/No-Go Decision Framework

GO criteria for mobile development:

- Database fully implemented and tested
- Authentication system functional
- Booking APIs working end-to-end
- Security vulnerabilities resolved
- At least 70% of MVP acceptance criteria met

Current Status:  **NO-GO** - Only 25-30% of MVP requirements met

CONCLUSION

The current Elara implementation represents a **promising UI/UX prototype** with functional AI integration, but it is **not production-ready** by any objective measure. When compared against the comprehensive Product/Tech Brief requirements, the application is missing:

- **91% of required database schema**
- **75% of required API endpoints**
- **100% of required microservices**
- **85% of core user stories**
- **100% of automation features**
- **All compliance infrastructure**

Summary Scorecard

Category	Score	Status
Architecture	5/100	 Critical
Database Schema	8/100	 Critical
API Endpoints	15/100	 Critical

Category	Score	Status
Microservices	0/100	X Critical
State Machine Logic	15/100	X Critical
User Stories	13/100	X Critical
Compliance	10/100	X Critical
Security	15/100	● Critical
OVERALL MVP READINESS	10/100	X NOT READY

Recommended Path Forward

1. Treat current implementation as Phase 0: UI/UX Prototype
2. Allocate 10-13 weeks for proper MVP development
3. Hire/assign 2 full-stack developers + 1 QA engineer
4. Follow the 3-phase roadmap outlined in Section 12
5. Re-assess after Phase 1 completion (database + auth)

Bottom Line

This codebase needs **substantial backend development** before it can be considered a professional, production-ready application suitable for mobile deployment. The UI/UX foundation is solid, but the infrastructure, data persistence, security, and automation layers are either missing or inadequate for production use.

Estimated Time to Production-Ready MVP: 10-13 weeks

Estimated Budget: xxxx (2-3 developers at market rates)

Recommended Next Step: Begin Phase 1(Database + Authentication) immediately

Report Compiled By: RevAI Agent

Assessment Date: October 10, 2025

Based On: Product/Tech Brief MVP Requirements

Next Review: After Phase 1 completion (estimated 4 weeks)

APPENDIX A: Current File Structure

```

elara/
└── client/
    └── src/

```

```

|   |   └── components/      # UI components (working)
|   |   └── pages/          # Route pages (working)
|   |   └── data/           # Static doctor data (10 doctors)
|   |   └── utils/          # Classifier, memory utils (partial)
|   |   └── lib/             # Query client setup
|   └── main.tsx
└── server/
    ├── index.ts            # Express server + OpenAI
    ├── routes.ts           # Health + mock endpoints
    ├── storage.ts          # In-memory storage (demo only)
    └── vite.ts              # Dev server integration
└── shared/
    └── schema.ts           # Drizzle schema (1 table only)
└── package.json

```

APPENDIX B: Required vs. Actual Endpoints

Required (8 endpoints):

1. POST /chat/send
2. GET /doctors/search
3. POST /bookings
4. POST /checkins
5. GET /mirror/weekly
6. GET /wearables/summary
7. POST /automations/schedule
8. GET /legal/disclaimer

Actual (5 endpoints):

1. POST /api/classify(partial)
2. POST /api/elara/reply(partial)
3. GET /api/mock/doctors(demo)
4. GET /api/mock/tests(demo)
5. GET /health(basic)

Gap: 3 of 8 required endpoints exist (37.5%), and only partially functional

End of Report